

Lab - 3

Name : Arul Kumar ARK

Roll No. : 225229103

Fuel Amount Prediction using Linear Regression

Step : 1

Prepare Your Dataset

Step : 2

```
In [2]: import pandas as pd
```

```
In [3]: fuel=pd.read_csv('fuel_data.csv')
fuel
```

Out[3]:

	drivenKM	fuelAmount
0	390.00	3600.0
1	403.00	3705.0
2	396.50	3471.0
3	383.50	3250.5
4	321.10	3263.7
5	391.30	3445.2
6	386.10	3679.0
7	371.80	3744.5
8	404.30	3809.0
9	392.20	3905.0
10	386.43	3874.0
11	395.20	3910.0
12	381.00	4020.7
13	372.00	3622.0
14	397.00	3450.5
15	407.00	4179.0
16	372.40	3454.2
17	375.60	3883.8
18	399.00	4235.9

```
In [6]: fuel.head(10)
```

Out[6]:

	drivenKM	fuelAmount
0	390.0	3600.0
1	403.0	3705.0
2	396.5	3471.0
3	383.5	3250.5
4	321.1	3263.7
5	391.3	3445.2
6	386.1	3679.0
7	371.8	3744.5
8	404.3	3809.0
9	392.2	3905.0

```
In [7]: fuel.shape
```

```
Out[7]: (19, 2)
```

```
In [8]: fuel.columns
```

```
Out[8]: Index(['drivenKM', 'fuelAmount'], dtype='object')
```

```
In [9]: type(fuel)
```

```
Out[9]: pandas.core.frame.DataFrame
```

```
In [10]: fuel.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19 entries, 0 to 18
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   drivenKM    19 non-null     float64
1   fuelAmount  19 non-null     float64
dtypes: float64(2)
memory usage: 432.0 bytes
```

Step : 3


```
In [11]: fuel.isnull()
```

```
Out[11]:
```

	drivenKM	fuelAmount
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
5	False	False
6	False	False
7	False	False
8	False	False
9	False	False
10	False	False
11	False	False
12	False	False
13	False	False
14	False	False
15	False	False
16	False	False
17	False	False
18	False	False

Step : 4

```
In [25]: import matplotlib.pyplot as plt
fuel.plot(kind="scatter",x="drivenKM",y="fuelAmount")
print(plt.show())
```



None

Step : 5 & Step : 6

```
In [15]: X = fuel[["drivenKM"]]
```

```
In [16]: print(X)
type(X)
```

```
drivenKM
0    390.00
1    403.00
2    396.50
3    383.50
4    321.10
5    391.30
6    386.10
7    371.80
8    404.30
9    392.20
10   386.43
11   395.20
12   381.00
13   372.00
14   397.00
15   407.00
16   372.40
17   375.60
18   399.00
```

Out[16]: pandas.core.frame.DataFrame

```
In [17]: y = fuel.fuelAmount
print(y)
type(y)
```

```
0    3600.0
1    3705.0
2    3471.0
3    3250.5
4    3263.7
5    3445.2
6    3679.0
7    3744.5
8    3809.0
9    3905.0
10   3874.0
11   3910.0
12   4020.7
13   3622.0
14   3450.5
15   4179.0
16   3454.2
17   3883.8
18   4235.9
Name: fuelAmount, dtype: float64
```

Out[17]: pandas.core.series.Series

Step : 7

```
In [19]: from sklearn.linear_model import LinearRegression
```

```
In [20]: from sklearn.model_selection import train_test_split
```

```
In [34]: X_train,X_test,y_train,y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [35]: X_train.shape
```

```
Out[35]: (15, 1)
```

```
In [36]: X_test.shape
```

```
Out[36]: (4, 1)
```

```
In [37]: y_train.shape
```

```
Out[37]: (15,)
```

```
In [38]: y_test.shape
```

```
Out[38]: (4,)
```

Part - I Linear Regression Model

Step : 8

```
In [39]: reg=LinearRegression()  
reg.fit(X_train,y_train)
```

```
Out[39]: LinearRegression()
```

Step : 9

```
In [40]: pred_800_KM=reg.predict([[800]])  
print("Deisel price for 800KM:",pred_800_KM)
```

```
Deisel price for 800KM: [6905.64571567]
```

```
C:\Users\arulk\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but Linear  
Regression was fitted with feature names  
warnings.warn(
```

Step : 10

```
In [41]: y_pred=reg.predict(X_test)  
y_pred
```

```
Out[41]: array([3775.81615646, 3785.74000628, 3815.51155575, 3875.05465468])
```

Step : 11

```
In [42]: import sklearn.metrics as metrics  
MSE=metrics.mean_squared_error(y_test,y_pred)  
R2=metrics.r2_score(y_test,y_pred)  
print("MSE: ",MSE.astype('int'))
```

```
MSE: 46181
```

```
In [33]: print("R2: ",R2)
```

```
R2: -0.4409983890088389
```

```
In [43]: print("coefficient:",reg.coef_)
```

```
coefficient: [7.63373063]
```

```
In [44]: ▶ print("Intercept:", reg.intercept_)
```

```
Intercept: 798.6612098962887
```

Part - II Linear Regression Model With Scaling Using StanderScaler

Step : 12 & Step : 13

```
In [50]: ▶ from sklearn.preprocessing import StandardScaler  
s=StandardScaler()  
sd_X_train=s.fit_transform(X_train)  
sd_X_train
```

```
Out[50]: array([[ 1.0601947 ],  
                [-0.5322439 ],  
                [ 0.02186483],  
                [-0.55221178],  
                [ 1.19497791],  
                [-0.37250084],  
                [ 0.670821  ],  
                [ 0.45616627],  
                [ 0.79562026],  
                [-3.09312478],  
                [-0.10293443],  
                [-0.56219572],  
                [ 0.16812957],  
                [ 0.69578085],  
                [ 0.15165606]])
```

```
In [51]: ▶ sd_X_test=s.transform(X_test)  
sd_X_test
```

```
Out[51]: array([[0.34634292],  
                [0.41123853],  
                [0.60592538],  
                [0.99529908]])
```

```
In [52]: ▶ reg.fit(sd_X_train,y_train)
```

```
Out[52]: LinearRegression()
```

```
In [53]: ▶ sd_y_pred=reg.predict(sd_X_test)  
sd_y_pred
```

```
Out[53]: array([3775.81615646, 3785.74000628, 3815.51155575, 3875.05465468])
```

Step : 14

```
In [54]: ▶ sd_MSE=metrics.mean_squared_error(y_test,sd_y_pred)  
sd_R2=metrics.r2_score(y_test,sd_y_pred)
```

```
In [55]: ▶ print("Mean Squared Error: ",sd_MSE)
```


```
Mean Squared Error: 46181.36710639172
```

```
In [56]: ▶ print("R2 Error: ",sd_R2)
```

```
R2 Error: -0.6180990161577082
```

Step : 15

```
In [57]: ▶ plt.scatter(y_test,y_pred,color='Brown',marker='o')
plt.title("Scatter plot between actual y vs predicted y values")
plt.xlabel('Actual y value')
plt.ylabel('Predicted y value')
plt.show()
```



Part - III Linear Regression Model With Scaling Using StanderScaler Using MinMaxScaler and Comparison with KNeighborsRegressor and SCDRegressor

Step : 16

```
In [29]: ▶ from sklearn.preprocessing import MinMaxScaler
mm=MinMaxScaler()
MinMax_X_train=mm.fit_transform(X_train)
MinMax_X_test=mm.transform(X_test)

reg.fit(MinMax_X_train,y_train)
MinMax_y_pred=reg.predict(MinMax_X_test)
print("Predictions of MinMaxScaler:",MinMax_y_pred)

MinMax_MSE=metrics.mean_squared_error(y_test,MinMax_y_pred)
MinMax_R2=metrics.r2_score(y_test,MinMax_y_pred)
print("MinMaxScaler MSE: ",MinMax_MSE)
print("MinMaxScaler R2: ",MinMax_R2)

Predictions of MinMaxScaler: [3775.81615646 3785.74000628 3815.51155575 3875.05465468]
MinMaxScaler MSE: 46181.3671063917
MinMaxScaler R2: -0.6180990161577073
```

Step : 17

```
In [30]: ▶ from sklearn.neighbors import KNeighborsRegressor
knr=KNeighborsRegressor()
knr.fit(X_train,y_train)
knr_y_pred=knr.predict(X_test)
print("Predictions of KNeighborsRegressor:",knr_y_pred)
knr_mse=metrics.mean_squared_error(y_test,knr_y_pred)
knr_r2=metrics.r2_score(y_test,knr_y_pred)
print("KNR MSE: ",knr_mse)
print("KNR R2: ",knr_r2)

Predictions of KNeighborsRegressor: [3635.9 3675.9 3787.28 3829.08]
KNR MSE: 21241.836200000045
KNR R2: 0.2557302563733307
```

Step : 18

```
In [31]: from sklearn.linear_model import SGDRegressor
sgd=SGDRegressor()
sgd.fit(X_train, y_train)
sgd_y_pred=sgd.predict(X_test)
print("Predictions of SGDRegressor:",sgd_y_pred)
sgd_mse=metrics.mean_squared_error(y_test, sgd_y_pred)
sgd_r2=metrics.r2_score(y_test,sgd_y_pred)
print("SGD MSE:",sgd_mse)
print("SGD R2:",sgd_r2)
```

Predictions of SGDRegressor: [-2.21183835e+14 -2.21921118e+14 -2.24132969e+14 -2.28556671e+14]
 SGD MSE: 5.016125282148493e+28
 SGD R2: -1.7575459308663273e+24

C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\linear_model\stochastic_gradient.py:128: FutureWarning: max_iter and tol parameters have been added in <class 'sklearn.linear_model.stochastic_gradient.SGDRegressor'> in 0.19. If both are left unset, they default to max_iter=5 and tol=None. If tol is not None, max_iter defaults to max_iter=1000. From 0.21, default max_iter will be 1000, and default tol will be 1e-3.
 "and default tol will be 1e-3." % type(self), FutureWarning)

Step : 19

```
In [32]: data_mse = {'lr_mse':[46181.36710639155], 'std_mse':[46181.36710639172], 'minmax_mse':[46181.3671063917], 'knr_mse':[21241.83620]

def best_model(data_mse):
    # Calculating the Lowest MSE
    mse_min = min(data_mse.values())
    # Storing the Lowest MSE in result
    result = [key for key in data_mse if data_mse[key] == mse_min]
    Model_name = []
    if result == ['lr_mse']:
        a = 'LinearRegression'
        Model_name.append(a)
    elif result == ['std_mse']:
        b = 'StandardScaler'
        Model_name.append(b)
    elif result == ['minmax_mse']:
        c = 'MinMaxScaler'
        Model_name.append(c)
    elif result == ['knr_mse']:
        d = 'KNeighborsRegressor'
        Model_name.append(d)
    elif result == ['sgd_mse']:
        e = 'SGDRegressor'
        Model_name.append(e)
    # Printing the result
    print("The best model with the lowest MSE to be selected is", Model_name)

best_model(data_mse)
```

The best model with the lowest MSE to be selected is ['KNeighborsRegressor']