

Name : Arul Kumar ARK

Roll No. : 225229103

Lab - 5

1 Diabetes Classification using Logistic Regression

Step : 1In [2]: 1 `import pandas as pd`In [3]: 1 `data = pd.read_csv("diabetes.csv")`In [4]: 1 `data.head(10)`

Out[4]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
5	5	116	74	0	0	25.6	0.201	30	0
6	3	78	50	32	88	31.0	0.248	26	1
7	10	115	0	0	0	35.3	0.134	29	0
8	2	197	70	45	543	30.5	0.158	53	1
9	8	125	96	0	0	0.0	0.232	54	1

In [5]: 1 `print(data.columns)`

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

In [6]: 1 `print(data.shape)`

(768, 9)

In [7]: 1 `print(data.dtypes)`

```
Pregnancies      int64
Glucose           int64
BloodPressure     int64
SkinThickness     int64
Insulin           int64
BMI               float64
DiabetesPedigreeFunction float64
Age               int64
Outcome           int64
dtype: object
```

In [8]: `print(data.info)`

```
<bound method DataFrame.info of
0      6      148      72      35      0      33.6
1      1      85      66      29      0      26.6
2      8     183      64      0      0      23.3
3      1      89      66      23      94      28.1
4      0     137      40      35     168      43.1
..    ...    ...    ...    ...    ...    ...
763    10     101      76      48     180      32.9
764     2     122      70      27      0      36.8
765     5     121      72      23     112      26.2
766     1     126      60      0      0      30.1
767     1      93      70      31      0      30.4

DiabetesPedigreeFunction  Age  Outcome
0      0.627      50      1
1      0.351      31      0
2      0.672      32      1
3      0.167      21      0
4      2.288      33      1
..    ...    ...    ...
763    0.171      63      0
764    0.340      27      0
765    0.245      30      0
766    0.349      47      1
767    0.315      23      0
```

[768 rows x 9 columns]>

In [13]: `print(data.value_counts)`

```
<bound method DataFrame.value_counts of
0      6      148      72      35      0      33.6
1      1      85      66      29      0      26.6
2      8     183      64      0      0      23.3
3      1      89      66      23      94      28.1
4      0     137      40      35     168      43.1
..    ...    ...    ...    ...    ...    ...
763    10     101      76      48     180      32.9
764     2     122      70      27      0      36.8
765     5     121      72      23     112      26.2
766     1     126      60      0      0      30.1
767     1      93      70      31      0      30.4

DiabetesPedigreeFunction  Age  Outcome
0      0.627      50      1
1      0.351      31      0
2      0.672      32      1
3      0.167      21      0
4      2.288      33      1
..    ...    ...    ...
763    0.171      63      0
764    0.340      27      0
765    0.245      30      0
766    0.349      47      1
767    0.315      23      0
```

[768 rows x 9 columns]>

step : 2

In [14]: `X = data.drop(['Outcome'],axis=1)`

In [15]:

1 X

Out[15]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33
...
763	10	101	76	48	180	32.9	0.171	63
764	2	122	70	27	0	36.8	0.340	27
765	5	121	72	23	112	26.2	0.245	30
766	1	126	60	0	0	30.1	0.349	47
767	1	93	70	31	0	30.4	0.315	23

768 rows × 8 columns

In [16]: 1 y = data['Outcome']

In [17]: 1 print(y)

```

0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64

```

```

In [23]: 1 from sklearn.model_selection import StratifiedShuffleSplit
2 splitter=StratifiedShuffleSplit(n_splits=1,test_size=0.25,random_state=12)
3 for train,test in splitter.split(X,y):
4     X_train = X.loc[train]
5     y_train = y.loc[train]
6     X_test = X.loc[test]
7     y_test = y.loc[test]

```

```

In [27]: 1 from sklearn.linear_model import LogisticRegression
2 model=LogisticRegression()
3 model.fit(X_train,y_train)
4 y_pred_m=model.predict(X_test)

```

C:\Users\arul\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n_iter_i = _check_optimize_result(

In [29]: 1 y_pred_m

```

Out[29]: array([1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0,
1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0,
0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0,
0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0,
0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1,
0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1], dtype=int64)

```

Step : 3

```
In [31]: 1 new=model.predict([[6,200,90,10,25,23.3,.672,42]])
2 if new==0:
3     print("Non-diabetic :",new)
4 else:
5     print("Diabetic : ",new)
```

Diabetic : [1]

C:\Users\arul\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
warnings.warn(

```
In [72]: 1 print('Accuracy : ',sum(y_test==y_pred_m)/float(y_test.shape[0]))
```

Accuracy : 0.796875

```
In [73]: 1 from sklearn.metrics import precision_score
2 print('Precision : ',precision_score(y_test,y_pred_m))
```

Precision : 0.7258064516129032

```
In [74]: 1 from sklearn.metrics import recall_score
2 print('Recall : ',recall_score(y_test,y_pred_m))
```

Recall : 0.6716417910447762

```
In [75]: 1 from sklearn.metrics import roc_auc_score
2 print('AUC Scoure : ',roc_auc_score(y_test,y_pred_m))
```

AUC Scoure : 0.767820895522388


Step : 4

```
In [76]: 1 from sklearn.metrics import confusion_matrix
2 matrix=confusion_matrix(y_test,y_pred_m)
3 matrix
```

Out[76]: array([[108, 17],
[22, 45]], dtype=int64)

```
In [77]: 1 import seaborn as sns
2 sns.heatmap(matrix, annot=True)
```

Out[77]: <AxesSubplot:>



Step : 5

```
In [78]: 1 from sklearn.preprocessing import MinMaxScaler
2 mm=MinMaxScaler()
3 mm_X_train=mm.fit_transform(X_train)
4 mm_X_train
```

```
Out[78]: array([[0.          , 0.62311558, 0.57377049, ..., 0.40834575, 0.07514944,
0.25          ],
[0.17647059, 0.96984925, 0.57377049, ..., 0.52011923, 0.06959863,
0.06666667],
[0.58823529, 0.61306533, 0.63934426, ..., 0.41132638, 0.1853117 ,
0.4          ],
...,
[0.29411765, 0.73869347, 0.63934426, ..., 0.50223547, 0.05977797,
0.73333333],
[0.17647059, 0.64824121, 0.52459016, ..., 0.39344262, 0.06020495,
0.11666667],
[0.          , 0.63316583, 0.70491803, ..., 0.40834575, 0.18659266,
0.          ]])
```

```
In [79]: 1 mm_X_test=mm.transform(X_test)
2 mm_X_test
```

```
Out[79]: array([[0.23529412, 0.77386935, 0.59016393, ..., 0.46646796, 0.11101623,
0.26666667],
[0.05882353, 0.78894472, 0.59016393, ..., 0.38152012, 0.01921435,
0.05          ],
[0.76470588, 0.79396985, 0.93442623, ..., 0.63040238, 0.0764304 ,
0.38333333],
...,
[0.05882353, 0.61306533, 0.73770492, ..., 0.74068554, 0.10546541,
0.16666667],
[0.17647059, 0.75376884, 0.62295082, ..., 0.31296572, 0.05508113,
0.26666667],
[0.11764706, 0.64321608, 0.52459016, ..., 0.59612519, 0.43680615,
0.05          ]])
```

```
In [80]: 1 mm_lor=LogisticRegression()
2 mm_lor=mm_lor.fit(mm_X_train,y_train)
3 mm_y_pred=mm_lor.predict(mm_X_test)
4 mm_y_pred
```

```
Out[80]: array([0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0,
0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0,
0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1,
0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0], dtype=int64)
```

```
In [81]: 1 print('Accuracy : ',sum(y_test==mm_y_pred)/float(y_test.shape[0]))
2
```

Accuracy : 0.7864583333333334

```
In [82]: 1 from sklearn.metrics import precision_score
2 print('Precision : ',precision_score(mm_y_pred,y_pred_m))
```

Precision : 0.8064516129032258

```
In [83]: 1 from sklearn.metrics import recall_score
2 print('Recall : ',recall_score(mm_y_pred,y_pred_m))
```

Recall : 1.0

```
In [84]: 1 from sklearn.metrics import roc_auc_score
2 print('AUC Scoure : ',roc_auc_score(mm_y_pred,y_pred_m))
```

AUC Scoure : 0.9577464788732395

Step : 6

```
In [85]: 1 from sklearn.preprocessing import StandardScaler
2 ss=StandardScaler()
3 ss_X_train=ss.fit_transform(X_train)
4 ss_X_train
```

```
Out[85]: array([[ -1.12598572,  0.11956974,  0.06558538, ..., -0.59413103,
        -0.65316149,  0.23834618],
       [ -0.23468782,  2.31457195,  0.06558538, ...,  0.34990533,
        -0.69286658, -0.70256434],
       [  1.84500729,  0.05594649,  0.46376382, ..., -0.56895673,
         0.13483187,  1.00818206],
       ...,
       [  0.35951079,  0.85123714,  0.46376382, ...,  0.19885951,
        -0.76311405,  2.71892846],
       [ -0.23468782,  0.27862787, -0.23304845, ..., -0.72000254,
        -0.76005981, -0.44595238],
       [ -1.12598572,  0.18319299,  0.86194226, ..., -0.59413103,
         0.14399459, -1.04471362]])
```

```
In [86]: 1 ss_X_test=ss.transform(X_test)
2 ss_X_test
```

```
Out[86]: array([[ 0.06241149,  1.07391853,  0.16512999, ..., -0.10323212,
        -0.39660551,  0.3238835 ],
       [ -0.82888642,  1.1693534 ,  0.16512999, ..., -0.82069975,
        -1.05326665, -0.78810166],
       [  2.73630519,  1.20116503,  2.2555668 , ...,  1.28135453,
        -0.64399878,  0.92264474],
       ...,
       [ -0.82888642,  0.05594649,  1.06103148, ...,  2.21280374,
        -0.4363106 , -0.18934042],
       [ -0.23468782,  0.94667202,  0.36421921, ..., -1.39970872,
        -0.79671067,  0.3238835 ],
       [ -0.53178712,  0.24681625, -0.23304845, ...,  0.99185005,
         1.93377797, -0.78810166]])
```

```
In [87]: 1 ss_lor=LogisticRegression()
2 ss_lor.fit(ss_X_train,y_train)
3 ss_y_pred=ss_lor.predict(ss_X_test)
```

```
In [88]: 1 print('Accuracy : ',sum(y_test==ss_y_pred)/float(y_test.shape[0]))
2
3 from sklearn.metrics import precision_score
4 print('Precision : ',precision_score(ss_y_pred,y_pred_m))
5
6 from sklearn.metrics import recall_score
7 print('Recall : ',recall_score(ss_y_pred,y_pred_m))
8
9 from sklearn.metrics import roc_auc_score
10 print('AUC Scoure : ',roc_auc_score(ss_y_pred,y_pred_m))
```

```
Accuracy : 0.7864583333333334
Precision : 0.967741935483871
Recall : 1.0
AUC Scoure : 0.9924242424242424
```


Step : 7

```
In [89]: 1 pred_prob1=mm_lor.predict_proba(mm_X_test)
```

```
In [90]: 1 from sklearn.metrics import roc_curve
2 fpr, tpr, thresh = roc_curve(y_test, pred_prob1[:,1], pos_label=1)
```

```
In [91]: 1 import matplotlib.pyplot as plt
2 plt.plot(fpr, tpr, linestyle='--', color='red', label='MinMaxScaler values')
```

Out[91]: [<matplotlib.lines.Line2D at 0x1d79c33be80>]



Step : 8

```
In [94]: 1 from sklearn.neighbors import KNeighborsClassifier
2 knn=KNeighborsClassifier(n_neighbors=4)
3 knn=knn.fit(X_train,y_train)
```

```
In [101]: 1 knn_y_pred=knn.predict(X_test)
```

```
In [102]: 1 from sklearn.preprocessing import MinMaxScaler
2 m=MinMaxScaler()
3 m_X_train=m.fit_transform(X_train)
4 m_X_train
```

Out[102]: array([[0. , 0.62311558, 0.57377049, ..., 0.40834575, 0.07514944,
0.25],
[0.17647059, 0.96984925, 0.57377049, ..., 0.52011923, 0.06959863,
0.06666667],
[0.58823529, 0.61306533, 0.63934426, ..., 0.41132638, 0.1853117 ,
0.4],
...,
[0.29411765, 0.73869347, 0.63934426, ..., 0.50223547, 0.05977797,
0.73333333],
[0.17647059, 0.64824121, 0.52459016, ..., 0.39344262, 0.06020495,
0.11666667],
[0. , 0.63316583, 0.70491803, ..., 0.40834575, 0.18659266,
0.]])

```
In [103]: 1 m_X_test=m.transform(X_test)
2 m_X_test
```

Out[103]: array([[0.23529412, 0.77386935, 0.59016393, ..., 0.46646796, 0.11101623,
0.26666667],
[0.05882353, 0.78894472, 0.59016393, ..., 0.38152012, 0.01921435,
0.05],
[0.76470588, 0.79396985, 0.93442623, ..., 0.63040238, 0.0764304 ,
0.38333333],
...,
[0.05882353, 0.61306533, 0.73770492, ..., 0.74068554, 0.10546541,
0.16666667],
[0.17647059, 0.75376884, 0.62295082, ..., 0.31296572, 0.05508113,
0.26666667],
[0.11764706, 0.64321608, 0.52459016, ..., 0.59612519, 0.43680615,
0.05]])

```
In [104]: 1 m_knn=KNeighborsClassifier()
2 m_knn=m_knn.fit(m_X_train,y_train)
```

```
In [105]: 1 m_y_pred=m_knn.predict(m_X_test)
          2 m_y_pred
```

```
Out[105]: array([0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0,
                0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0,
                1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
                0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0,
                1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0,
                0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0,
                0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1,
                0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1,
                0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1,
                0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int64)
```

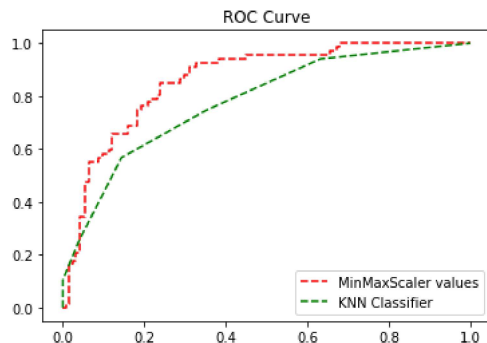
```
In [108]: 1 print('Accuracy : ',sum(y_test==m_y_pred)/float(y_test.shape[0]))
          2
          3 from sklearn.metrics import precision_score
          4 print('Precision : ',precision_score(m_y_pred,y_pred_m))
          5
          6 from sklearn.metrics import recall_score
          7 print('Recall : ',recall_score(m_y_pred,y_pred_m))
          8
          9 from sklearn.metrics import roc_auc_score
         10 print('AUC Scoure : ',roc_auc_score(m_y_pred,y_pred_m))
```

```
Accuracy : 0.7552083333333334
Precision : 0.6935483870967742
Recall : 0.7678571428571429
AUC Scoure : 0.8140756302521008
```

Step : 9

```
In [109]: 1 pred_prob2=m_knn.predict_proba(m_X_test)
```

```
In [110]: 1 from sklearn.metrics import roc_curve
          2 fpr1,tpr1,thresh1=roc_curve(y_test,pred_prob2[:,1],pos_label=1)
          3 import matplotlib.pyplot as plt
          4 plt.plot(fpr1,tpr1,linestyle='--',color='red',label='MinMaxScaler values')
          5 plt.plot(fpr1,tpr1,linestyle='--',color='green',label='KNN Classifier')
          6 plt.title('ROC Curve')
          7 plt.legend(loc='best')
          8 plt.savefig('ROC',dpi=300)
          9 plt.show()
```



Step : 10

```
In [114]: 1 from sklearn.linear_model import LogisticRegressionCV
          2 model1=LogisticRegressionCV(Cs=10,cv=4,penalty='l1',solver='liblinear')
          3 model2=LogisticRegressionCV(Cs=10,cv=4,penalty='l2')
          4 model1.fit(mm_X_train,y_train)
          5 model2.fit(mm_X_train,y_train)
```

```
Out[114]: LogisticRegressionCV(cv=4)
```

```
In [116]: 1 rg_y_pred1 = model1.predict(mm_X_test)
          2 rg_y_pred2 = model2.predict(mm_X_test)
```

AUC SCORE OF L1


```
In [118]: 1 from sklearn.metrics import roc_auc_score
2 l1auc = roc_auc_score(y_test, rg_y_pred1)
3 l1auc = (' LOR L1 MinMax AUC', l1auc)
4 print(l1auc)
```

```
(' LOR L1 MinMax AUC', 0.7465074626865671)
```


AUC SCORE OF L2

```
In [119]: 1 l2auc = roc_auc_score(y_test, rg_y_pred2)
2 l2auc = (' LOR L2 MinMax AUC', l2auc)
3 l2auc
```

```
Out[119]: (' LOR L2 MinMax AUC', 0.7568955223880597)
```

Step : 11

```
In [122]: 1 pred_prb7 = model1.predict_proba(mm_X_test)
2 pred_prb8 = model2.predict_proba(mm_X_test)
3 fpr,tbr,threshold = roc_curve(y_test, pred_prob1[:,1],pos_label=1)
4 fpr1,tbr1,threshold1 = roc_curve(y_test, pred_prob2[:,1],pos_label=1)
5 fpr2,tbr2,threshold2= roc_curve(y_test, pred_prb7[:,1],pos_label=1)
6 fpr3,tbr3,threshold3 = roc_curve(y_test, pred_prb8[:,1],pos_label=1)
7 plt.plot(fpr, tbr, linestyle='-', color='green', label='LogisticRegression')
8 plt.plot(fpr1, tbr1, linestyle='-', color='red', label='KNN')
9 plt.plot(fpr3, tbr3, linestyle='-', color='blue', label='l2')
10 plt.plot(fpr2, tbr2, linestyle='-', color='yellow', label='l1')
11 plt.title('Receiver Operating Characteristic')
12 plt.legend(loc = 'best')
13 plt.ylabel('True Positive Rate')
14 plt.xlabel('False Positive Rate')
15 plt.show()
```



```
In [ ]: 1
```